

SENTIMENT ANALYSIS OF TWEETS

A Project Report on
SENTIMENT ANALYSIS OF TWEETS

Submitted in partial fulfillment of the requirements for
the degree of

Bachelor of Technology

in

Computer Science Engineering

BY

KONDAL SRINIDHI SAGAR

```

# House Price Prediction (Supervised Learning - Regression)
# Predict home prices using ML models on California Housing Dataset

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
import pickle

# -----
# 1. Load and Explore the Data
# -----
print("Loading California housing data...")
housing = fetch_california_housing()
df = pd.DataFrame(housing.data, columns=housing.feature_names)
df['Price'] = housing.target

print("\nFirst few rows of data:")
print(df.head())

print(f"\nDataset shape: {df.shape}")
print(f"Price range: ${df['Price'].min()*100000:,.2f} - ${df['Price'].max()*100000:,.2f}")

# -----
# 2. Visualize the Data
# -----
print("\nVisualizing data...")
plt.figure(figsize=(12, 8))
df.hist(bins=30, figsize=(15, 10))
plt.suptitle("Feature Distributions", y=1.02)
plt.tight_layout()
plt.show()

plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Feature Correlation Matrix")
plt.show()

# -----
# 3. Feature Selection & Preprocessing
# -----
# Drop Latitude and Longitude for simplicity
X = df.drop(['Price', 'Latitude', 'Longitude'], axis=1)
y = df['Price']

# -----
# 4. Train/Test Split
# -----
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

```

```

# -----
# 5. Train Models
# -----
models = {
    "Linear Regression": LinearRegression(),
    "Random Forest": RandomForestRegressor(random_state=42),
    "XGBoost": XGBRegressor(random_state=42)
}

results = []

for name, model in models.items():
    print(f"\nTraining {name}...")
    model.fit(X_train, y_train)
    preds = model.predict(X_test)

    r2 = r2_score(y_test, preds)
    mae = mean_absolute_error(y_test, preds)
    rmse = np.sqrt(mean_squared_error(y_test, preds))

    results.append({
        "Model": name,
        "R2 Score": r2,
        "MAE ($)": mae * 100000,
        "RMSE ($)": rmse * 100000
    })

    print(f"R2: {r2:.3f}, MAE: ${mae*100000:,.2f}, RMSE: ${rmse*100000:,.2f}")

# -----
# 6. Model Evaluation
# -----
print("\nModel Evaluation Summary:")
results_df = pd.DataFrame(results)
print(results_df.sort_values("R2 Score", ascending=False))

# -----
# 7. Save Best Model
# -----
best_model = XGBRegressor(random_state=42)
best_model.fit(X_train, y_train)

with open("california_model.pkl", "wb") as f:
    pickle.dump(best_model, f)
print("\nBest model saved as 'california_model.pkl'")

# -----
# 8. Streamlit Web App (Optional)
# -----
def run_web_app():
    import streamlit as st

    st.title("California House Price Predictor")
    st.write("Estimate house price based on selected features.")

    st.sidebar.header("Input Features")
    MedInc = st.sidebar.slider("Median Income (in $10,000)", 0.0, 15.0, 3.0)
    HouseAge = st.sidebar.slider("House Age", 0, 100, 30)
    AveRooms = st.sidebar.slider("Average Rooms", 1.0, 15.0, 5.0)
    AveBedrms = st.sidebar.slider("Average Bedrooms", 0.5, 5.0, 1.0)

```

```

Population = st.sidebar.slider("Block Population", 0, 10000, 1500)
AveOccup = st.sidebar.slider("Average Occupancy", 1.0, 10.0, 3.0)

if st.sidebar.button("Predict Price"):
    input_features = [[MedInc, HouseAge, AveRooms, AveBedrms, Population, Av
with open("california_model.pkl", "rb") as f:
    model = pickle.load(f)
    prediction = model.predict(input_features)[0]
    st.success(f"Estimated House Price: ${prediction*100000:,.2f}")

# To run the web app:
# 1. Save this script as app.py
# 2. Run: streamlit run app.py
# run_web_app() # Uncomment to Launch

# -----
# End
# -----

print("\nYou can now:")
print("- Use the saved model for predictions")
print("- Run run_web_app() to launch a Streamlit app")

```

Loading California housing data...

First few rows of data:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude \
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85

	Longitude	Price
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422

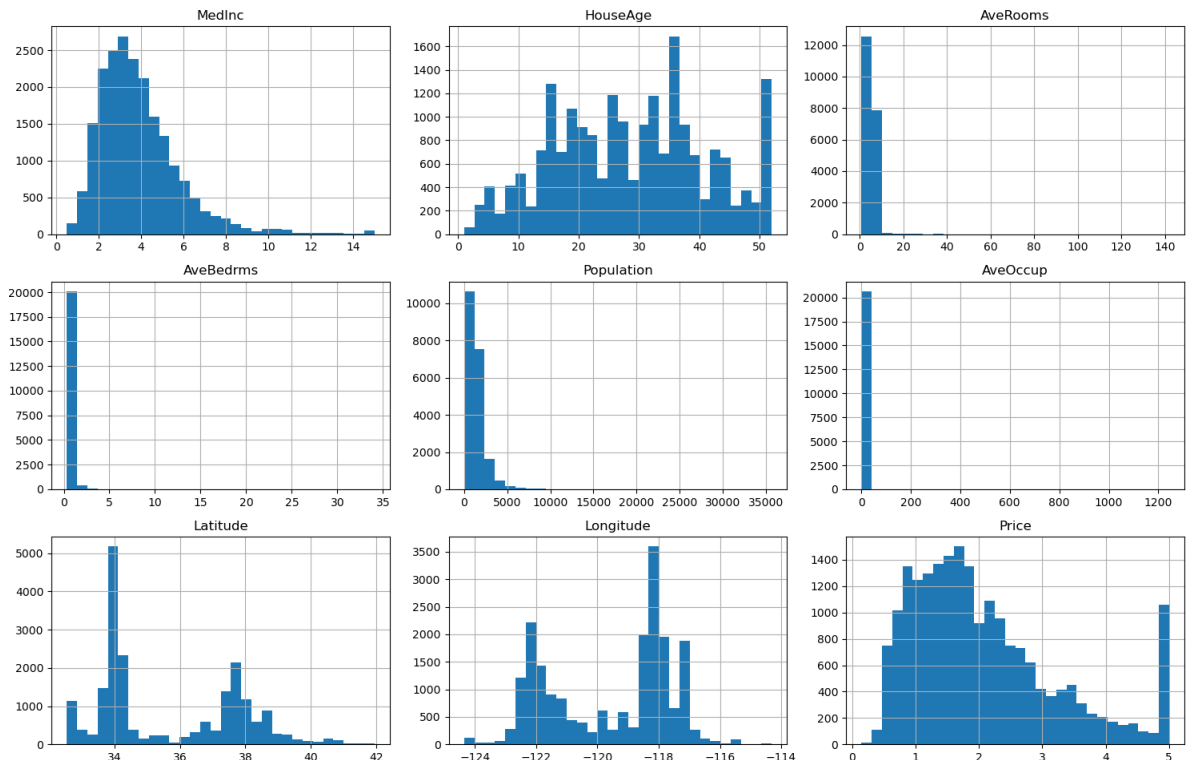
Dataset shape: (20640, 9)

Price range: \$14,999.00 - \$500,001.00


Visualizing data...

<Figure size 1200x800 with 0 Axes>

Feature Distributions



Feature Correlation Matrix



Training Linear Regression...

R²: 0.510, MAE: \$57,921.41, RMSE: \$80,136.59

Training Random Forest...

R²: 0.678, MAE: \$46,191.21, RMSE: \$65,005.54

Training XGBoost...

R²: 0.671, MAE: \$46,800.32, RMSE: \$65,639.16

Model Evaluation Summary:

	Model	R ² Score	MAE (\$)	RMSE (\$)
1	Random Forest	0.677527	46191.207871	65005.544472
2	XGBoost	0.671210	46800.316097	65639.155268
0	Linear Regression	0.509934	57921.406655	80136.585369

Best model saved as 'california_model.pkl'

You can now:

- Use the saved model for predictions
- Run `run_web_app()` to launch a Streamlit app


```

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_rep

# 1. Load the dataset
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['species'] = pd.Categorical.from_codes(iris.target, iris.target_names)

# 2. Vis  - Pairplot with Set1 palette
print("  Creating pairplot with Set1")
sns.pairplot(df, hue="species", palette="Set1")
plt.suptitle("Pairplot (Set1 Colors)", y=1.02)
plt.show()

# 3. Vis  - Boxplot with Pastel1 palette (with hue to prevent deprecation w
print("  Creating boxplot with Pastel1")
plt.figure(figsize=(10, 6))
sns.boxplot(x="species", y="petal length (cm)", data=df, hue="species", palette=
plt.title("Petal Length by Species (Pastel1)")
plt.show()

# 4. Vis  - Violin plot with Dark2 palette
print("  Creating violin plot with Dark2")
plt.figure(figsize=(10, 6))
sns.violinplot(x="species", y="sepal width (cm)", data=df, hue="species", palett
plt.title("Sepal Width by Species (Dark2)")
plt.show()

# 5. Prepare features and labels
X = df.drop("species", axis=1)
y = df["species"]

# 6. Split the data
X_train, X_test, y_train, y_test = train_test_split(

```

```

X, y, test_size=0.2, stratify=y, random_state=42
)

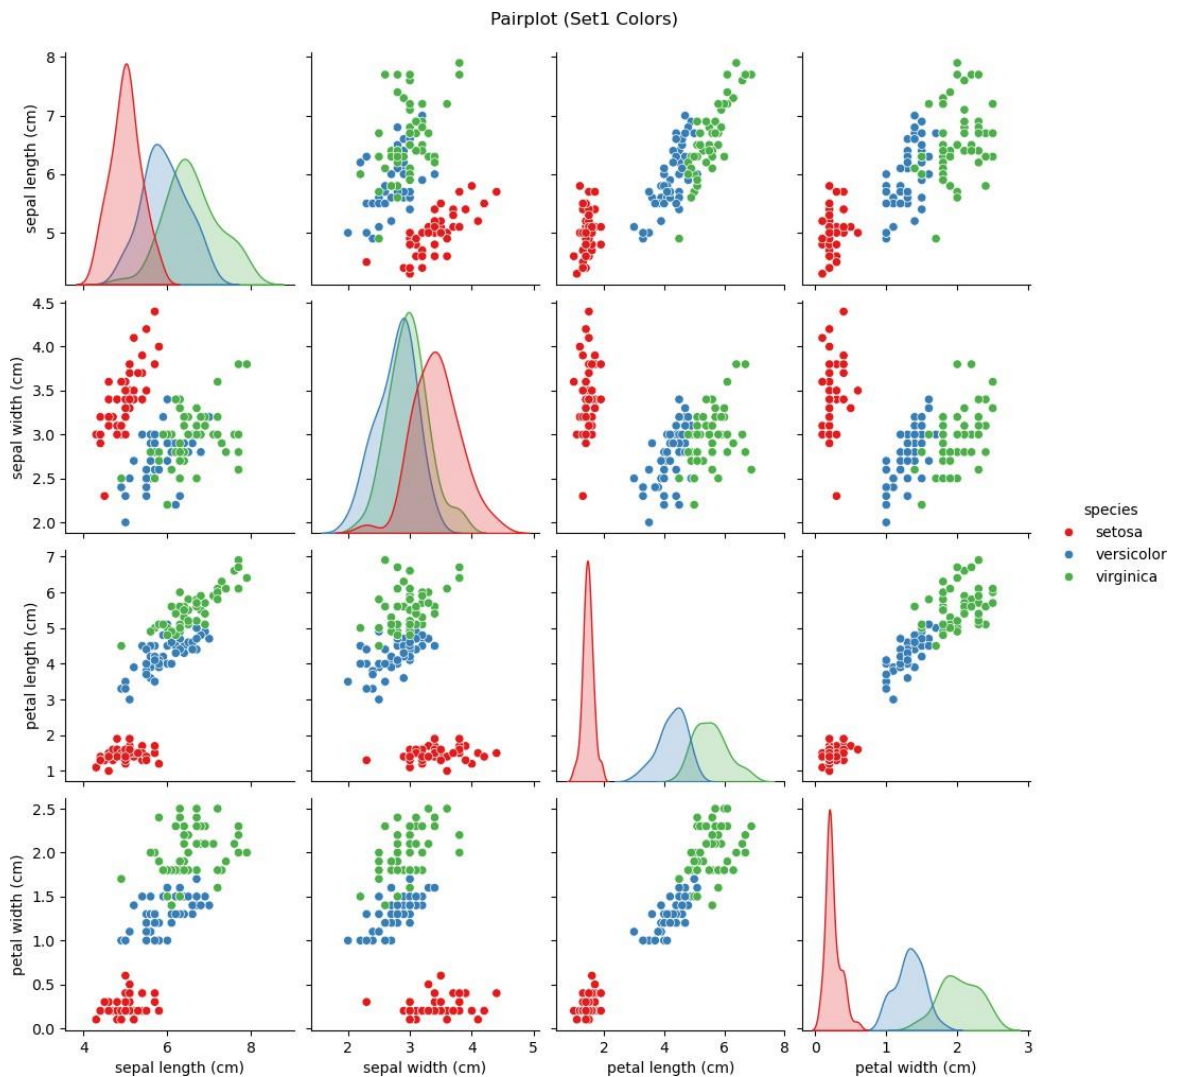
# 7. Train the model
model = KNeighborsClassifier(n_neighbors=3)
model.fit(X_train, y_train)

# 8. Evaluate the model
y_pred = model.predict(X_test)
acc = accuracy_score(y_test, y_pred)

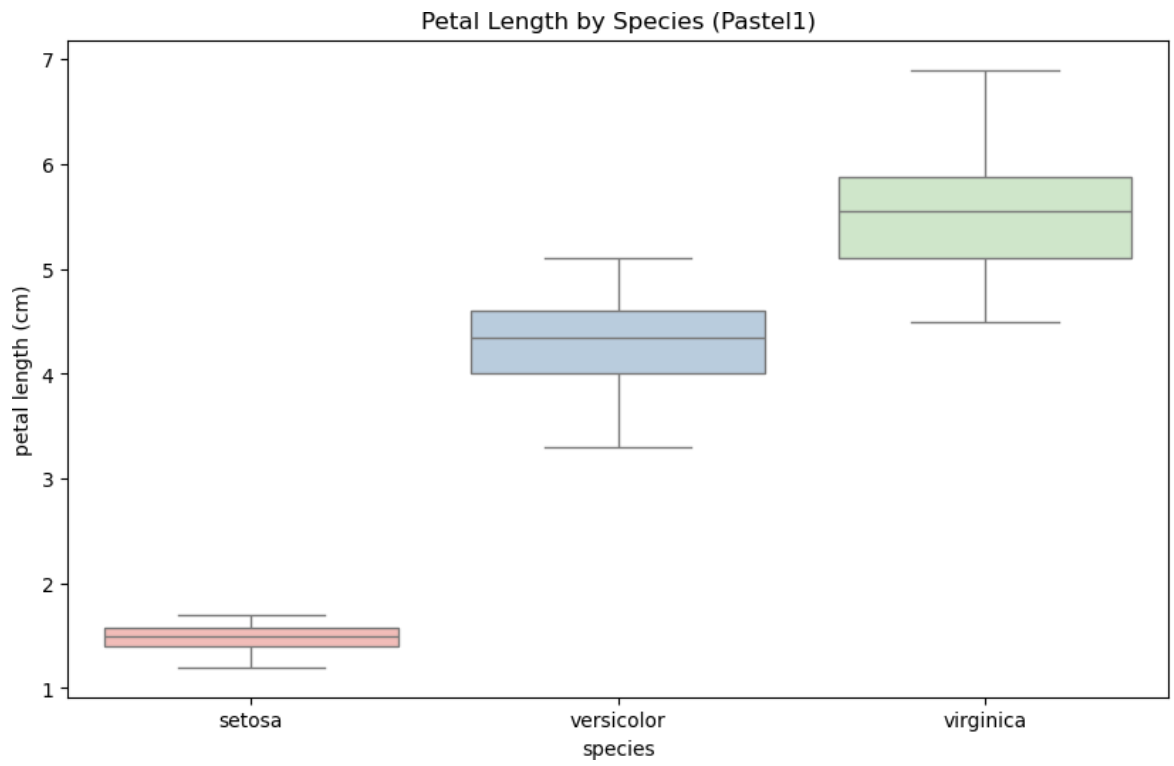
print(f"\n ✓ Accuracy: {acc:.2%}")
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```


Creating pairplot with Set1 palette...



Creating boxplot with Pastel1 palette...



Creating violin plot with Dark2 palette...





Accuracy: 100.00%

Confusion Matrix:

```
[[10  0  0]
 [ 0 10  0]
 [ 0  0 10]]
```

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	10
virginica	1.00	1.00	1.00	10
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30